

Plotting trees from Random Forest models with ggraph

- Tweet
-
-

16 March 2017

Today, I want to show how I use Thomas Lin Pedersen's awesome `ggraph` (<https://github.com/thomasp85/ggraph>) package to plot decision trees from Random Forest models.

I am very much a visual person, so I try to plot as much of my results as possible because it helps me get a better feel for what is going on with my data.

A nice aspect of using tree-based machine learning, like Random Forest models, is that that they are more easily interpreted than e.g. neural networks as they are based on decision trees. So, when I am using such models, I like to plot final decision trees (if they aren't too large) to get a sense of which decisions are underlying my predictions.

There are a few very convient ways to plot the outcome if you are using the `randomForest` package but I like to have as much control as possible about the layout, colors, labels, etc. And because I didn't find a solution I liked for `caret` models, I developed the following little function (below you may find information about how I built the model):

As input, it takes part of the output from `model_rf <- caret::train(... "rf" ...)`, that gives the trees of the final model: `model_rf$finalModel$forest`. From these trees, you can specify which one to plot by index.

```

library(dplyr)
library(ggraph)
library(igraph)

tree_func <- function(final_model,
                      tree_num) {

  # get tree by index
  tree <- randomForest::getTree(final_model,
                                k = tree_num,
                                labelVar = TRUE) %>%
    tibble::rownames_to_column() %>%
    # make leaf split points to NA, so the 0s won't get plotted
    mutate(`split point` = ifelse(is.na(prediction), `split point`, NA))

  # prepare data frame for graph
  graph_frame <- data.frame(from = rep(tree$rowname, 2),
                            to = c(tree$`left daughter`, tree$`right daughter`))

  # convert to graph and delete the last node that we don't want to plot
  graph <- graph_from_data_frame(graph_frame) %>%
    delete_vertices("0")

  # set node labels
  V(graph)$node_label <- gsub("_", " ", as.character(tree$`split var`))
  V(graph)$leaf_label <- as.character(tree$prediction)
  V(graph)$split <- as.character(round(tree$`split point`, digits = 2))

  # plot
  plot <- ggraph(graph, 'dendrogram') +
    theme_bw() +
    geom_edge_link() +
    geom_node_point() +
    geom_node_text(aes(label = node_label), na.rm = TRUE, repel = TRUE) +
    geom_node_label(aes(label = split), vjust = 2.5, na.rm = TRUE, fill = "white") +
    geom_node_label(aes(label = leaf_label, fill = leaf_label), na.rm = TRUE,
                    repel = TRUE, colour = "white", fontface = "bold", show.le
gend = FALSE) +
    theme(panel.grid.minor = element_blank(),
          panel.grid.major = element_blank(),
          panel.background = element_blank(),
          plot.background = element_rect(fill = "white"),
          panel.border = element_blank(),
          axis.line = element_blank(),
          axis.text.x = element_blank(),
          axis.text.y = element_blank(),
          axis.ticks = element_blank(),
          axis.title.x = element_blank(),
          axis.title.y = element_blank(),
          plot.title = element_text(size = 18))

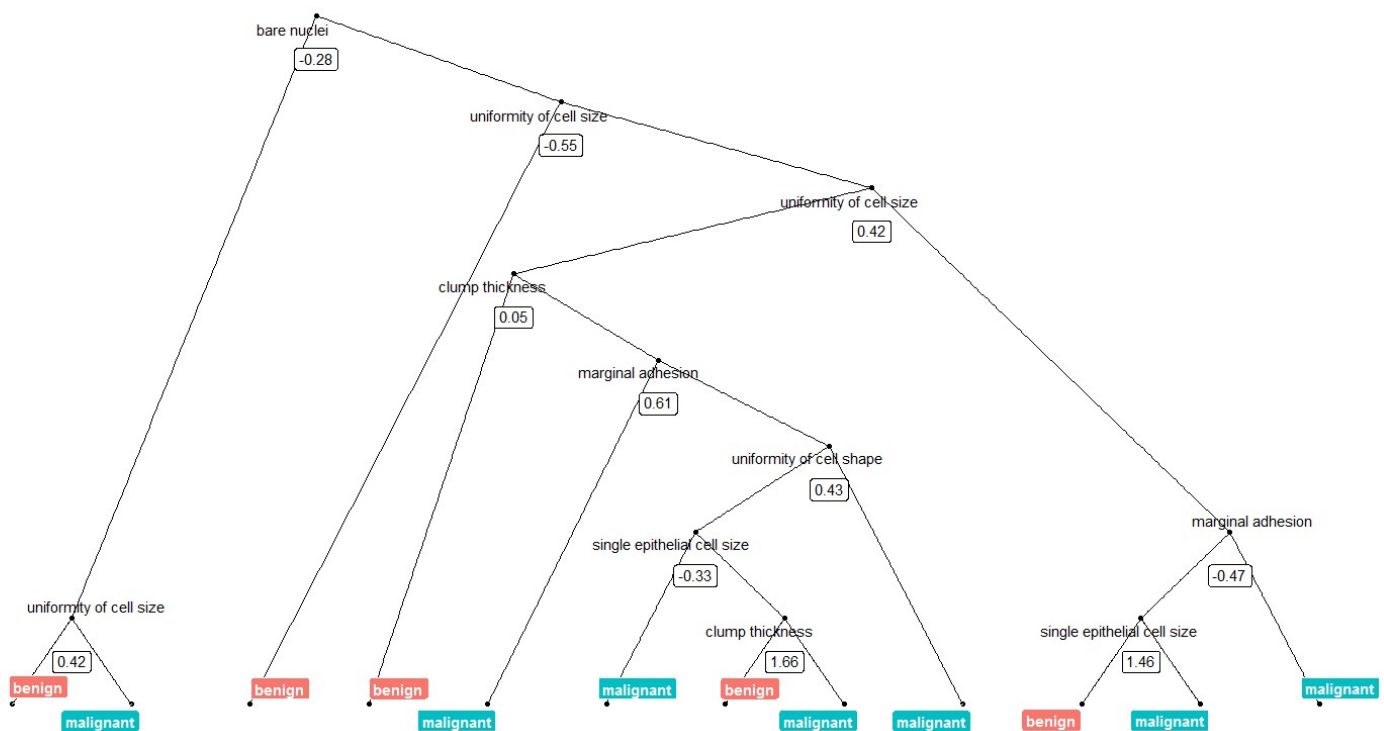
```

```
print(plot)
}
```

We can now plot, e.g. the tree with the smallest number of nodes:

```
tree_num <- which(model_rf$finalModel$forest$ndbigtree == min(model_rf$finalModel$forest$ndbigtree))

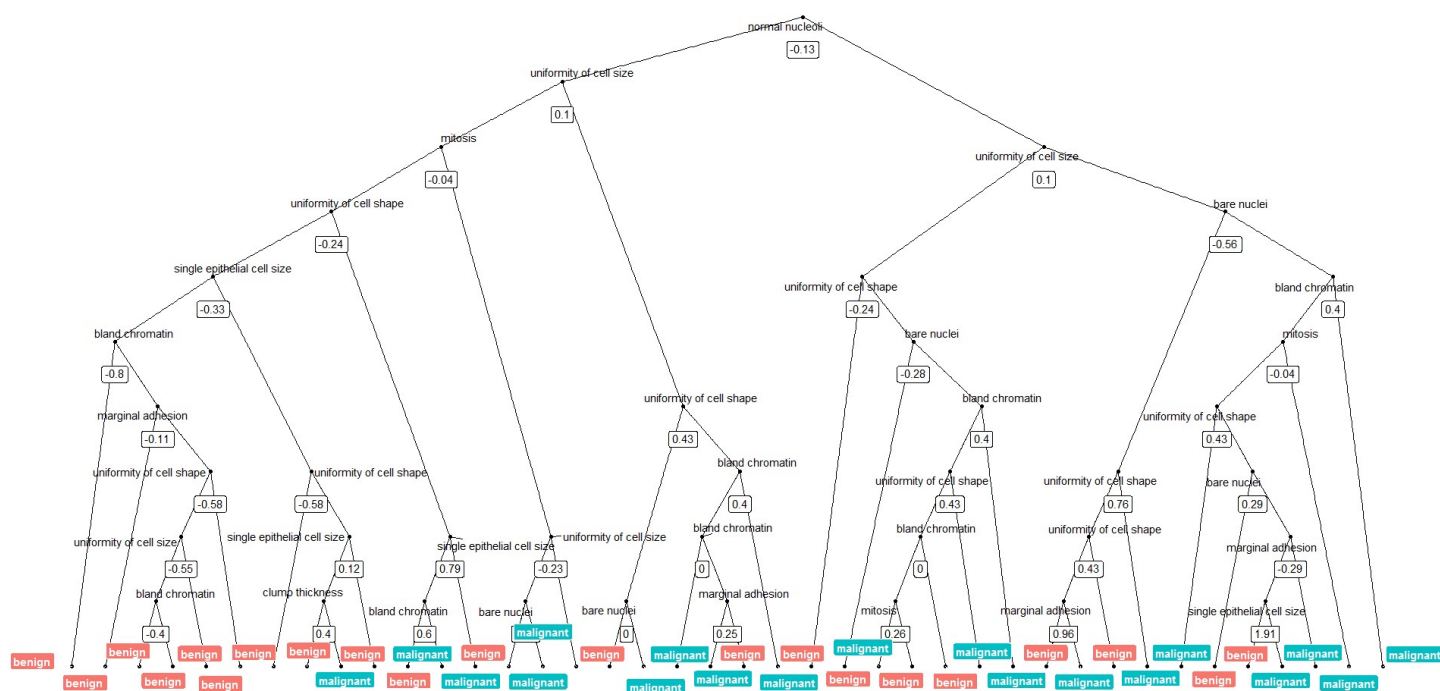
tree_func(final_model = model_rf$finalModel, tree_num)
```



Or we can plot the tree with the biggest number of nodes:

```
tree_num <- which(model_rf$finalModel$forest$ndbigtree == max(model_rf$finalModel$forest$ndbigtree))

tree_func(final_model = model_rf$finalModel, tree_num)
```



Preparing the data and modeling

The data set I am using in these example analyses, is the **Breast Cancer Wisconsin (Diagnostic) Dataset**. The data was downloaded from the UC Irvine Machine Learning Repository (<http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Diagnostic%29>).

The first data set looks at the predictor classes:

- malignant or
- benign breast mass.

The features characterize cell nucleus properties and were generated from image analysis of fine needle aspirates (FNA) (https://en.wikipedia.org/wiki/Fine-needle_aspiration) of breast masses:

- Sample ID (code number)
- Clump thickness
- Uniformity of cell size
- Uniformity of cell shape
- Marginal adhesion
- Single epithelial cell size
- Number of bare nuclei
- Bland chromatin
- Number of normal nuclei
- Mitosis
- Classes, i.e. diagnosis

```
bc_data <- read.table("datasets/breast-cancer-wisconsin.data.txt", header = FALSE, sep = ",")
colnames(bc_data) <- c("sample_code_number",
                      "clump_thickness",
                      "uniformity_of_cell_size",
                      "uniformity_of_cell_shape",
                      "marginal_adhesion",
                      "single_epithelial_cell_size",
                      "bare_nuclei",
                      "bland_chromatin",
                      "normal_nucleoli",
                      "mitosis",
                      "classes")

bc_data$classes <- ifelse(bc_data$classes == "2", "benign",
                        ifelse(bc_data$classes == "4", "malignant", NA))

bc_data[bc_data == "?"] <- NA

# impute missing data
library(mice)

bc_data[,2:10] <- apply(bc_data[, 2:10], 2, function(x) as.numeric(as.character(x)))
dataset_impute <- mice(bc_data[, 2:10], print = FALSE)
bc_data <- cbind(bc_data[, 11, drop = FALSE], mice::complete(dataset_impute, 1))

bc_data$classes <- as.factor(bc_data$classes)

# how many benign and malignant cases are there?
summary(bc_data$classes)

# separate into training and test data
library(caret)

set.seed(42)
index <- createDataPartition(bc_data$classes, p = 0.7, list = FALSE)
train_data <- bc_data[index, ]
test_data <- bc_data[-index, ]

# run model
set.seed(42)
model_rf <- caret::train(classes ~ .,
                        data = train_data,
                        method = "rf",
                        preProcess = c("scale", "center"),
                        trControl = trainControl(method = "repeatedcv",
                                                number = 10,
                                                repeats = 10,
                                                savePredictions = TRUE,
                                                verboseIter = FALSE))
```

If you are interested in more machine learning posts, check out the category listing for **machine_learning** on my blog (https://shiring.github.io/categories.html#machine_learning-ref).

```
sessionInfo()
```

```
## R version 3.3.3 (2017-03-06)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
## Running under: Windows 7 x64 (build 7601) Service Pack 1
##
## locale:
## [1] LC_COLLATE=English_United States.1252
## [2] LC_CTYPE=English_United States.1252
## [3] LC_MONETARY=English_United States.1252
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United States.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods    base
##
## other attached packages:
## [1] igraph_1.0.1      ggraph_1.0.0      ggplot2_2.2.1.9000
## [4] dplyr_0.5.0
##
## loaded via a namespace (and not attached):
## [1] Rcpp_0.12.9      nloptr_1.0.4      plyr_1.8.4
## [4] viridis_0.3.4    iterators_1.0.8    tools_3.3.3
## [7] digest_0.6.12    lme4_1.1-12        evaluate_0.10
## [10] tibble_1.2        gtable_0.2.0       nlme_3.1-131
## [13] lattice_0.20-34   mgcv_1.8-17        Matrix_1.2-8
## [16] foreach_1.4.3     DBI_0.6            ggrepel_0.6.5
## [19] yaml_2.1.14       parallel_3.3.3     SparseM_1.76
## [22] gridExtra_2.2.1    stringr_1.2.0      knitr_1.15.1
## [25] MatrixModels_0.4-1 stats4_3.3.3        rprojroot_1.2
## [28] grid_3.3.3         caret_6.0-73        nnet_7.3-12
## [31] R6_2.2.0           rmarkdown_1.3       minqa_1.2.4
## [34] udunits2_0.13      tweenr_0.1.5        deldir_0.1-12
## [37] reshape2_1.4.2     car_2.1-4           magrittr_1.5
## [40] units_0.4-2        backports_1.0.5     scales_0.4.1
## [43] codetools_0.2-15   ModelMetrics_1.1.0  htmltools_0.3.5
## [46] MASS_7.3-45        splines_3.3.3       randomForest_4.6-12
## [49] assertthat_0.1     pbkrtest_0.4-6      ggforce_0.1.1
## [52] colorspace_1.3-2   labeling_0.3         quantreg_5.29
## [55] stringi_1.1.2      lazyeval_0.2.0      munsell_0.4.3
```



[machine_learning¹⁴ \(/categories.html#machine_learning-ref\)](https://shiring.github.io/categories.html#machine_learning-ref)

[ggplot2³² \(/tags.html#ggplot2-ref\)](/tags.html#ggplot2-ref)[machine_learning¹³ \(/tags.html#machine_learning-ref\)](/tags.html#machine_learning-ref)[caret³ \(/tags.html#caret-ref\)](/tags.html#caret-ref)[random_forest⁶ \(/tags.html#random_forest-ref\)](/tags.html#random_forest-ref)[« Previous \(/machine_learning/2017/03/07/grid_search/\)](/machine_learning/2017/03/07/grid_search/)[Archive \(/archive.html\)](/archive.html)[Next » \(/machine_learning/2017/03/31/webinar_code/\)](/machine_learning/2017/03/31/webinar_code/)

Sponsored

Jedź łyżkę tego przed snem a na zawsze pozbędziesz się bólu stawów i całego ciała

Health Institute

A gdybyś rok temu zainwestował tysiąc dolarów w Teslę?

eToro

The Japanese Way To Remove Body Toxins

Detox Patches

Kobieta złożyła pozew o rozwód po zobaczeniu tego zdjęcia. Czy widzisz dlaczego?

PDFWonder

Poszukujemy 500 testerów aparatów słuchowych w Polsce

Ekspert Słuchu

Przeszkadza Ci częstomocz i prostata? Ekspert wyjaśnia prostą metodę.

Urostrom

Małopolskie: 80% refundacji na naukę języków. Rekrutacja trwa...

40 Comments

shirinsplayground

Disqus' Privacy Policy

1 Login ▾

Recommend 7

Tweet

Share

Sort by Best ▾



Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS (?)



Name

**Pieter Vos** • 4 years ago

Hi Shirin,

Thank you for sharing this. I have a question though. If you have 500 trees in your random forest, do you visualize each of them to understand your model? How would you know this single decision tree is representative for the random forest, because the final decision is made on all trees (ensemble)?

thanks

Pieter

2 ^ | ▾ • Reply • Share ›

**Shirin Elsinghorst** Mod ➔ Pieter Vos • 4 years ago

Hi Pieter.

Thanks! :-)

No, I don't look at all trees but I generally pick a few to plot so that I get an idea of which decisions were made by the model. Really understanding a reasonably complex model is by definition almost impossible but by looking at a few trees you can get a good sense of which features are similar and how they were split.

^ | ▾ • Reply • Share ›

**AJT** • 2 years ago

Brilliant post. Thanks & super cool!

1 ^ | ▾ • Reply • Share ›

**รักษ์ อย่างไรเงื่อนไข** • 3 years ago

Thank you so much, It great to share the random forest it showed in the decision trees.

1 ^ | ▾ • Reply • Share ›

**Glorious1** • 4 years ago • edited

Thank you for this. I have huge trees, with about 15,000 nodes (which I guess become vertices?). So I'm very interested in how you used `delete_vertices()`. I don't quite follow the comment " # . . . delete the last node that we don't want to plot".

1. Does this mean it should plot everything after the vertex you specify, but not before?

2. This doesn't seem to match the help for `delete_vertices`, which indicates you specify a location

and error when I try to plot (Error in seq.default(h[1], h[2], length.out = n) :
'to' must be a finite number).

Thanks for any help.

1 ^ | v • Reply • Share ›



Shirin Elsinghorst Mod → Glorious1 • 4 years ago

I am not plotting the last row of vertices ("0") because I plotted the labels at their stead.
As for you error messages, I can't really say without having seen your data and models...

^ | v • Reply • Share ›



This comment was deleted.



Shirin Elsinghorst Mod → Guest • 4 years ago

This is shown for the randomForest package...


^ | v • Reply • Share ›



GOKHAN CAMLIYURT • 2 months ago

Hi Shirin,

Thank oyu very much for sharing. Able to do a beatiful tree with little mistakes :)





22.08.2021, 16:33

^ | v • Reply • Share ›



This comment was marked as spam.



Shirin Elsinghorst Mod → JVD • 3 years ago

At which point in the function do you get this error?

^ | v • Reply • Share ›



JVD → Shirin Elsinghorst • 3 years ago • edited

EDIT:

I started everything again and somehow the error disappeared...

^ | v • Reply • Share ›



Shirin Elsinghorst Mod → JVD • 3 years ago

It's magic :-D

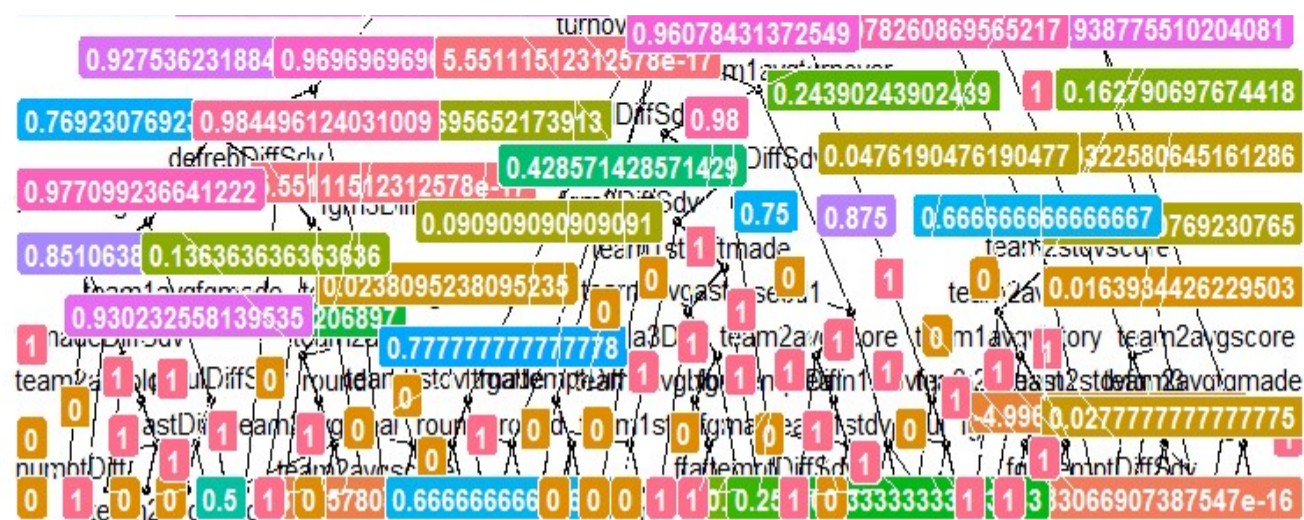
1 ^ | v • Reply • Share ›



Siladitya Sen • 4 years ago

Great! It worked!!!

^ | v • Reply • Share ›



 Thank you for the suggestion! I'll look into it. :)

1 ^ | v • Reply • Share ›



Abhijit Dasgupta • 4 years ago

You might want to mention early in your code that you are using random forests via the caret train

Sponsored

Jedź łyżkę tego przed snem a na zawsze pozbędziesz się bólu stawów i całego ciała

Health Institute

A gdybyś rok temu zainwestował tysiąc dolarów w Teslę?

eToro

Poszukujemy 500 testerów aparatów słuchowych w Polsce

Ekspert Słuchu

The Japanese Way To Remove Body Toxins

Detox Patches

Kobieta złożyła pozew o rozwód po zobaczeniu tego zdjęcia. Czy widzisz dlaczego?

PDFWonder

Jeżeli masz ponad 30 lat i posiadasz komputer, to musisz zagrać w tę grę!

Total Battle - Online Strategy Game



/Shirin_Glander) (<http://de.linkedin.com/in/shirin-glander-01120881>) with help from Jekyll



Bootstrap (<http://jekyllbootstrap.com>) and Bootstrap (<http://getbootstrap.com>)